



ScalaStan

Joe Wingbermuehle



Scala

"Object-Oriented Meets Functional"

2017 RedMonk Ranking:

- 1 JavaScript
- 2 Java
- 3 Python
- 4 PHP
- 5 C#
- 6 C++
- 7 CSS
- 8 Ruby
- 9 C
- 10 Objective-C
- 11 Swift
- 12 Shell
- **12 Scala**
- 14 R
- 15 Go
- 15 Perl

A hybrid object-oriented and functional programming language with a strong static type system.


Runs on the Java Virtual Machine (JVM).

```
object HelloWorld extends App {  
    println("Hello, World!")  
}
```

<https://scala-lang.org>


Goals of ScalaStan

- **Make it easy and type-safe to get data into Stan.**




Goals of ScalaStan

- **Make it easy and type-safe to get data into Stan.**
- **Make it easy and type-safe to get results out of Stan.**



Goals of ScalaStan

- Make it easy and type-safe to get data into Stan.
- Make it easy and type-safe to get results out of Stan.
- **Make running Stan simple.**



Goals of ScalaStan

- **Make it easy and type-safe to get data into Stan.**
- **Make it easy and type-safe to get results out of Stan.**
- **Make running Stan simple.**
- **Make developing Stan models easier:**
 - **Type safety**




Goals of ScalaStan

- Make it easy and type-safe to get data into Stan.
- Make it easy and type-safe to get results out of Stan.
- Make running Stan simple.
- **Make developing Stan models easier:**

- Type safety



- **Model and result caching**



Goals of ScalaStan

- Make it easy and type-safe to get data into Stan.
- Make it easy and type-safe to get results out of Stan.
- Make running Stan simple.
- **Make developing Stan models easier:**

- Type safety



- Model and result caching



- Discoverability of Stan functions and distributions



Goals of ScalaStan

- Make it easy and type-safe to get data into Stan.
- Make it easy and type-safe to get results out of Stan.
- Make running Stan simple.
- **Make developing Stan models easier:**

- Type safety



- Model and result caching



- Discoverability of Stan functions and distributions



- **Make Stan models composable**



Goals of ScalaStan

- Make it easy and type-safe to get data into Stan.
- Make it easy and type-safe to get results out of Stan.
- Make running Stan simple.
- **Make developing Stan models easier:**

- Type safety



- Model and result caching



- Discoverability of Stan functions and distributions



- Make Stan models composable



- **Make it possible to generate Stan**



What Is ScalaStan?

- A Stan wrapper in Scala



Like RStan and PyStan



- An embedded domain-specific language (DSL) for Stan
- Type-checked API in Scala for generating Stan



```
val model = new Model {  
  sigma ~ stan.cauchy(0, 1)  
  y ~ stan.normal(m * x + b, sigma)  
}
```



```
model {  
  sigma ~ cauchy(0, 1);  
  y ~ normal(m * x + b, sigma);  
}
```

Linear Regression Example

Stan Code:

```
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  real beta;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(alpha + beta * x, sigma);  
}
```

Simple linear regression:

$$y_n \sim \text{Normal}(\alpha + \beta X_n, \sigma)$$



Models in ScalaStan

Stan:

```
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}
```

ScalaStan:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

Models in ScalaStan

Stan:

```
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  real beta;  
  real<lower=0> sigma;  
}
```

ScalaStan:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))  
  
val alpha = parameter(real())  
val beta = parameter(real())  
val sigma = parameter(real(lower = 0))
```

Models in ScalaStan

Stan:

```
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real alpha;  
  real beta;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(alpha + beta * x, sigma);  
}
```

ScalaStan:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))  
  
val alpha = parameter(real())  
val beta = parameter(real())  
val sigma = parameter(real(lower = 0))  
  
val model = new Model {  
  | y ~ stan.normal(alpha + beta * x, sigma)  
}
```

Model is a trait in Scala
providing the DSL for
models.

Goal: Keep the low-level ScalaStan API as similar to Stan as possible.

Scala IDE Support

Types are checked in Scala

```
val model = new Model {  
  | y ~ stan.normal(alpha + beta * x, sigma) + sigma  
}
```

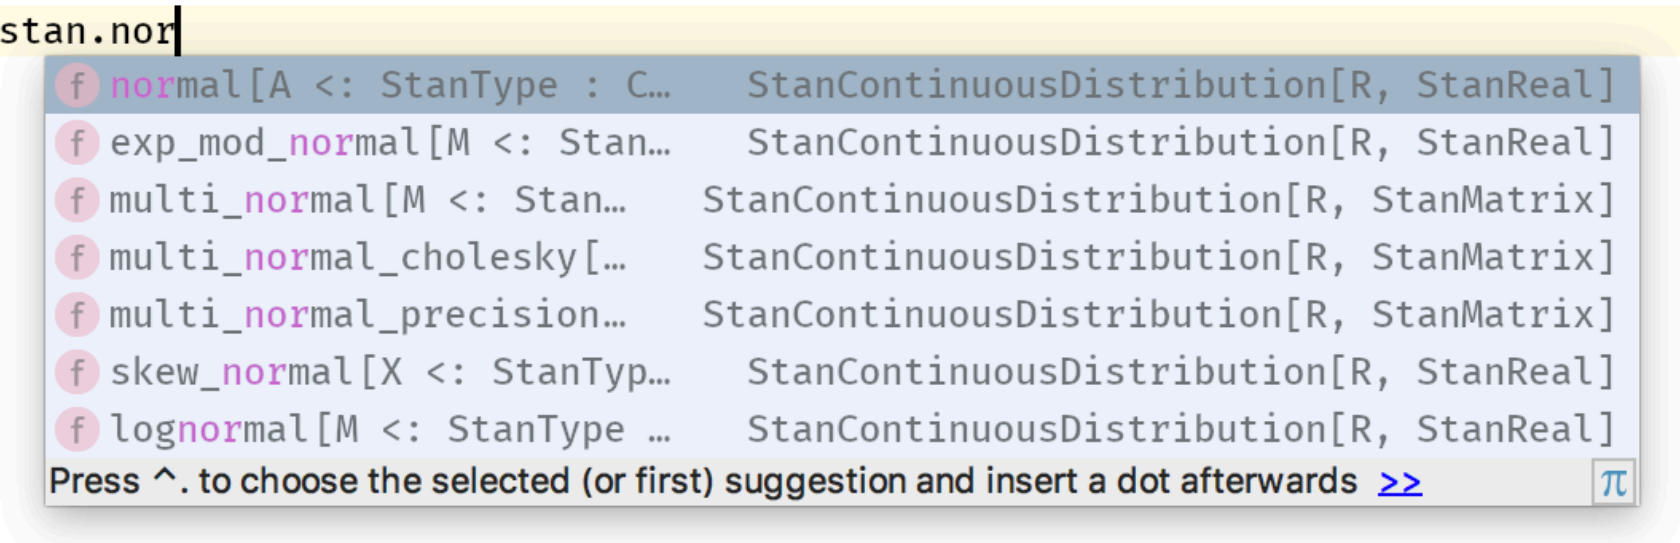

Scala IDE Support

Types are checked in Scala


```
val model = new Model {  
  y ~ stan.normal(alpha + beta * x, sigma) + sigma  
}
```

IDE provides discoverability

```
val model = new Model {  
  y ~ stan.nor  
}
```



- f normal[A <: StanType : C... StanContinuousDistribution[R, StanReal]
- f exp_mod_normal[M <: Stan... StanContinuousDistribution[R, StanReal]
- f multi_normal[M <: Stan... StanContinuousDistribution[R, StanMatrix]
- f multi_normal_cholesky[... StanContinuousDistribution[R, StanMatrix]
- f multi_normal_precision... StanContinuousDistribution[R, StanMatrix]
- f skew_normal[X <: StanTyp... StanContinuousDistribution[R, StanReal]
- f lognormal[M <: StanType ... StanContinuousDistribution[R, StanReal]

Press ^. to choose the selected (or first) suggestion and insert a dot afterwards >> 

Scala IDE Support


Types are checked in Scala

```
val model = new Model {  
  y ~ stan.normal(alpha + beta * x, sigma) + sigma  
}
```

IDE provides discoverability

```
val model = new Model {  
  y ~ stan.nor  
}
```

- f normal[A <: StanType : C... StanContinuousDistribution[R, StanReal]
- f exp_mod_normal[M <: Stan... StanContinuousDistribution[R, StanReal]
- f multi_normal[M <: Stan... StanContinuousDistribution[R, StanMatrix]
- f multi_normal_cholesky[... StanContinuousDistribution[R, StanMatrix]
- f multi_normal_precision... StanContinuousDistribution[R, StanMatrix]
- f skew_normal[X <: StanTyp... StanContinuousDistribution[R, StanReal]
- f lognormal[M <: StanType ... StanContinuousDistribution[R, StanReal]

Press ^. to choose the selected (or first) suggestion and insert a dot afterwards >> 

mu: StanValue[A], sigma: StanValue[B]

```
val model = new Model {  
  y ~ stan.normal(alpha + beta * x, sigma)  
}
```

Getting Data Into ScalaStan

Required data values:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

Getting Data Into ScalaStan

Required data values:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

From Scala directly:

```
val xs = Vector(1.0, 3.0, 4.0, 5.0)  
val ys = Vector(2.5, 5.1, 6.3, 7.0)  
val results = model  
  .withData(x, xs)  
  .withData(y, ys)
```



```
final def withData[T <: StanType, V](  
  decl: StanDataDeclaration[T],  
  data: V  
) (implicit ev: V <:: T#SCALA_TYPE): CompiledModel
```

Getting Data Into ScalaStan

Required data values:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

From Scala directly:

```
val xs = Vector(1.0, 3.0, 4.0, 5.0)  
val ys = Vector(2.5, 5.1, 6.3, 7.0)  
val results = model  
  .withData(x, xs)  
  .withData(y, ys)
```



```
final def withData[T <: StanType, V](  
  decl: StanDataDeclaration[T],  
  data: V  
) (implicit ev: V <:: T#SCALA_TYPE): CompiledModel
```

From CSV:

```
val source = CsvDataSource.fromFile("data.csv")  
val results = model  
  .withData(source(x, "x"))  
  .withData(source(y, "y"))
```

Getting Data Into ScalaStan

Required data values:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

From Scala directly:

```
val xs = Vector(1.0, 3.0, 4.0, 5.0)  
val ys = Vector(2.5, 5.1, 6.3, 7.0)  
val results = model  
  .withData(x, xs)  
  .withData(y, ys)
```



```
final def withData[T <: StanType, V](  
  decl: StanDataDeclaration[T],  
  data: V  
)(implicit ev: V <:: T#SCALA_TYPE): CompiledModel
```

From CSV:

```
val source = CsvDataSource.fromFile("data.csv")  
val results = model  
  .withData(source(x, "x"))  
  .withData(source(y, "y"))
```

From R:

```
val rsource = RDataSource.fromFile("data.R")
```

Getting Data Into ScalaStan

Required data values:

```
val N = data(int(lower = 0))  
val y = data(vector(N))  
val x = data(vector(N))
```

From Scala directly:

```
val xs = Vector(1.0, 3.0, 4.0, 5.0)  
val ys = Vector(2.5, 5.1, 6.3, 7.0)  
val results = model  
  .withData(x, xs)  
  .withData(y, ys)
```



```
final def withData[T <: StanType, V](  
  decl: StanDataDeclaration[T],  
  data: V  
)(implicit ev: V <:: T#SCALA_TYPE): CompiledModel
```

From CSV:

```
val source = CsvDataSource.fromFile("data.csv")  
val results = model  
  .withData(source(x, "x"))  
  .withData(source(y, "y"))
```

From R:

```
val rsource = RDataSource.fromFile("data.R")
```

- Scala ensures the types line up
- The value for N is set automatically

Running ScalaStan Models

```
val results = model  
  .withData(x, xs)  
  .withData(y, ys)  
  .run()
```


Running ScalaStan Models

```
val results = model  
  .withData(x, xs)  
  .withData(y, ys)  
  .run()
```

Multiple chains
(to get parallelism and
convergence testing):

```
.run(chains = 4)
```

Running ScalaStan Models

```
val results = model  
  .withData(x, xs)  
  .withData(y, ys)  
  .run()
```

Multiple chains
(to get parallelism and
convergence testing):

```
.run(chains = 4)
```

```
final def run(  
  chains: Int = 1,  
  seed: Int = -1,  
  cache: Boolean = true,  
  method: RunMethod.Method = RunMethod.Sample()  
): StanResults = {
```

Other options:

- Fix seed
- Result caching
- Sampler/optimizer parameters

Getting Results Out of ScalaStan

```
def mean[T <: StanType](decl: StanParameterDeclaration[T]): T#SUMMARY_TYPE
```

```
def sd[T <: StanType](decl: StanParameterDeclaration[T]): T#SUMMARY_TYPE
```

```
results.summary(System.out)
println(s"alpha = ${results.mean(alpha)}")
println(s"beta = ${results.mean(beta)}")
```




4 chains with 1000 iterations each, 4000 total

Name	Mean	MCSE	StdDev	5%	50%	95%	N_Eff	R_hat
accept_stat__	0.7195	0.03290	0.3413	0.005329	0.8987	0.9999	107.6	1.008
alpha	2.367	0.7259	8.851	-1.166	1.546	5.351	148.7	1.010
beta	0.8654	0.2706	2.748	-0.02792	1.134	1.933	103.1	1.013
divergent__	0.05925	0.008916	0.2361	0.000	0.000	1.000	701.4	1.001
energy__	2.495	0.6340	3.228	-1.479	1.851	8.476	25.93	1.081
lp__	-0.9925	0.7085	2.978	-6.615	-0.2887	2.454	17.67	1.091
n_leapfrog__	18.03	1.661	31.38	3.000	15.00	47.00	357.2	1.001
sigma	2.356	1.016	9.548	0.2312	0.7624	5.971	88.34	1.016
stepsize__	0.07296	0.03059	0.04327	0.04306	0.05132	0.1477	2.001	Infinity
treedepth__	3.215	0.09472	1.259	1.000	3.000	5.000	176.8	1.010

```
alpha = 2.367276227341493
beta = 0.865363928450001
```

Neal's Funnel

```
val x = parameter(real())  
val y = parameter(real())  
  
val model = new Model {  
  y ~ stan.normal(0, 3)  
  x ~ stan.normal(0, stan.exp(y / 2))  
}
```



72 of 1000 iterations ended up with a divergence

Non-Centered Parameterization

```
val rawx = parameter(real())
```


```
val rawy = parameter(real())
```

```
val y = new TransformedParameter(real()) {  
  | result := 3 * rawy  
}
```

```
val x = new TransformedParameter(real()) {  
  | result := stan.exp(y / 2) * rawx  
}
```

```
val model = new Model {  
  | y ~ stan.normal(0, 1)  
  | x ~ stan.normal(0, 1)  
}
```

No divergences!



Non-centered Parameterization

```
trait TransformedNormalSampler extends Model {  
  def sample(  
    value: ParameterDeclaration[StanReal],  
    mu: StanValue[StanReal],  
    sigma: StanValue[StanReal]  
  ): ParameterDeclaration[StanReal] = {  
  
    value ~ stan.normal(0, 1)  
  
    new TransformedParameter(real()) {  
      result := mu + sigma * value  
    }  
  }  
}
```

```
val model = new Model with TransformedNormalSampler {  
  val y = sample(rawy, 0.0, 3.0)  
  val x = sample(rawx, 0.0, stan.exp(y / 2))  
}
```

Composing Models using ScalaStan

Multiple Linear Regression using Least Squares:

Put the model in a trait and implement it with a case class:

```
trait LeastSquaresLike extends Model {  
  val x: DataDeclaration[StanMatrix]  
  val y: DataDeclaration[StanVector]  
  val beta: ParameterDeclaration[StanVector]  
  
  target += -stan.dotSelf(y - x * beta)  
}
```

```
case class LeastSquaresRegression(  
  x: DataDeclaration[StanMatrix],  
  y: DataDeclaration[StanVector],  
  beta: ParameterDeclaration[StanVector]  
) extends LeastSquaresLike
```

 Subtract the squared error from target

Composing Models using ScalaStan

Multiple Linear Regression using Least Squares:

Put the model in a trait and implement it with a case class:

```
trait LeastSquaresLike extends Model {  
  val x: DataDeclaration[StanMatrix]  
  val y: DataDeclaration[StanVector]  
  val beta: ParameterDeclaration[StanVector]  
  
  target += -stan.dotSelf(y - x * beta)  
}  
  
case class LeastSquaresRegression(  
  x: DataDeclaration[StanMatrix],  
  y: DataDeclaration[StanVector],  
  beta: ParameterDeclaration[StanVector]  
) extends LeastSquaresLike
```

Subtract the squared error from target

Now we can create the parameters/data values and instantiate the case class:

```
val n = data(int(lower = 0))  
val p = data(int(lower = 0))  
val y = data(vector(n))  
val x = data(matrix(n, p))  
  
val beta = parameter(vector(p))  
  
val model = LeastSquaresRegression(x, y, beta)
```


Composing Models using ScalaStan

Multiple Linear Regression with the Ridge Penalty:

Implement the ridge penalty in a trait by updating the log probability (target):

```
trait RidgePenaltyLike extends Model {  
  val ridgeLambda: DataDeclaration[StanReal]  
  val beta: ParameterDeclaration[StanVector]  
  
  target += -ridgeLambda * stan.dotSelf(beta)  
}
```

↑
Minimize the Euclidean length of the coefficients

```
case class RidgeRegression(  
  x: DataDeclaration[StanMatrix],  
  y: DataDeclaration[StanVector],  
  beta: ParameterDeclaration[StanVector],  
  ridgeLambda: DataDeclaration[StanReal]  
) extends LeastSquaresLike with RidgePenaltyLike
```

↑
Linear regression trait from before

Composing Models using ScalaStan

Multiple Linear Regression with the Lasso penalty:

```
trait LassoPenaltyLike extends Model {  
  val lassoLambda: DataDeclaration[StanReal]  
  val beta: ParameterDeclaration[StanVector]  
  
  for (i ← beta.range) {  
    target += -lassoLambda * stan.fabs(beta(i))  
  }  
}
```

Minimize the sum of the absolute values of the coefficients

```
case class LassoRegression(  
  x: DataDeclaration[StanMatrix],  
  y: DataDeclaration[StanVector],  
  beta: ParameterDeclaration[StanVector],  
  lassoLambda: DataDeclaration[StanReal]  
) extends LeastSquaresLike with LassoPenaltyLike
```

Composing Models using ScalaStan

Multiple Linear Regression with the Elastic Net:

```
case class ElasticNetRegression(  
  x: DataDeclaration[StanMatrix],  
  y: DataDeclaration[StanVector],  
  beta: ParameterDeclaration[StanVector],  
  ridgeLambda: DataDeclaration[StanReal],  
  lassoLambda: DataDeclaration[StanReal]  
) extends LeastSquaresLike with RidgePenaltyLike with LassoPenaltyLike {  
  val betaElasticNet = new GeneratedQuantity(vector(stan.cols(x))) {  
    result := (1 + ridgeLambda) * beta  
  }  
}
```

Both penalties



```
val model = ElasticNetRegression(x, y, beta, lambda1, lambda2)
```

Other Features

Locals and Assignment:

```
val x = local(real())  
x := 10
```

Other Features

Locals and Assignment:

```
val x = local(real())  
x := 10
```

Conditionals:

```
when(x > 10) {  
  // ...  
} otherwise {  
  // ...  
}
```

Other Features

Locals and Assignment:

```
val x = local(real())  
x := 10
```

Conditionals:

```
when(x > 10) {  
  // ...  
} otherwise {  
  // ...  
}
```

"while" loops:

```
loop(x < 10) {  
  // ...  
}
```

Other Features

Locals and Assignment:

```
val x = local(real())  
x := 10
```

Parameter (and Data) Transforms:

```
val beta2 = new TransformedParameter(real()) {  
  | result := beta * 2  
}
```

Conditionals:

```
when(x > 10) {  
  // ...  
} otherwise {  
  // ...  
}
```

"while" loops:

```
loop(x < 10) {  
  // ...  
}
```

Other Features

Locals and Assignment:

```
val x = local(real())  
x := 10
```

Conditionals:

```
when(x > 10) {  
    // ...  
} otherwise {  
    // ...  
}
```

"while" loops:

```
loop(x < 10) {  
    // ...  
}
```

Parameter (and Data) Transforms:

```
val beta2 = new TransformedParameter(real()) {  
    result := beta * 2  
}
```

Functions:

```
val func = new Function(real()) {  
    val x = input(real())  
    output(x + 1)  
}
```

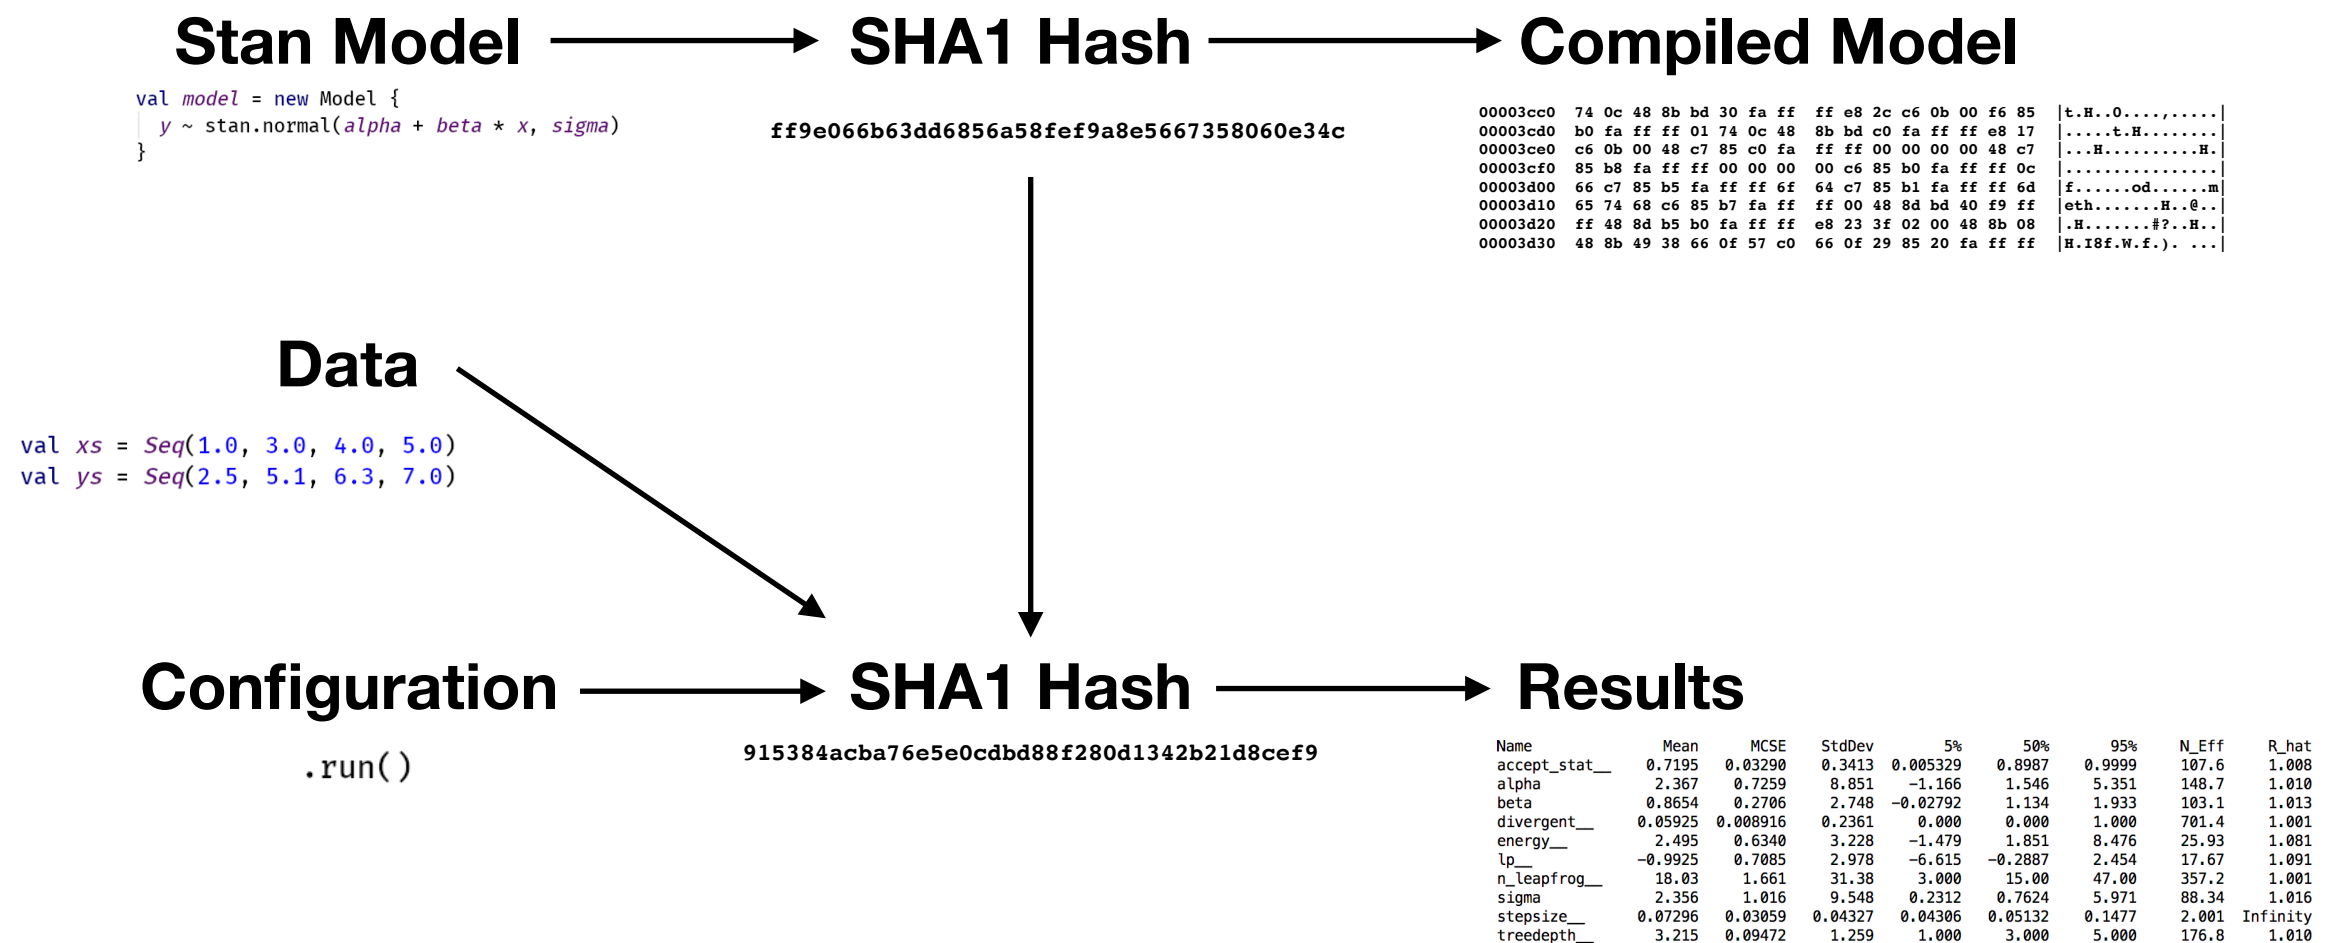
Return value

Parameter
declaration

Return type


Model/Result Caching

- Avoid re-compiling/building the model when the model doesn't change
- Avoid re-running the model when the model and data don't change
 - Make post-processing changes faster to test.




Integration into the Stan Ecosystem

Model Re-use: Export

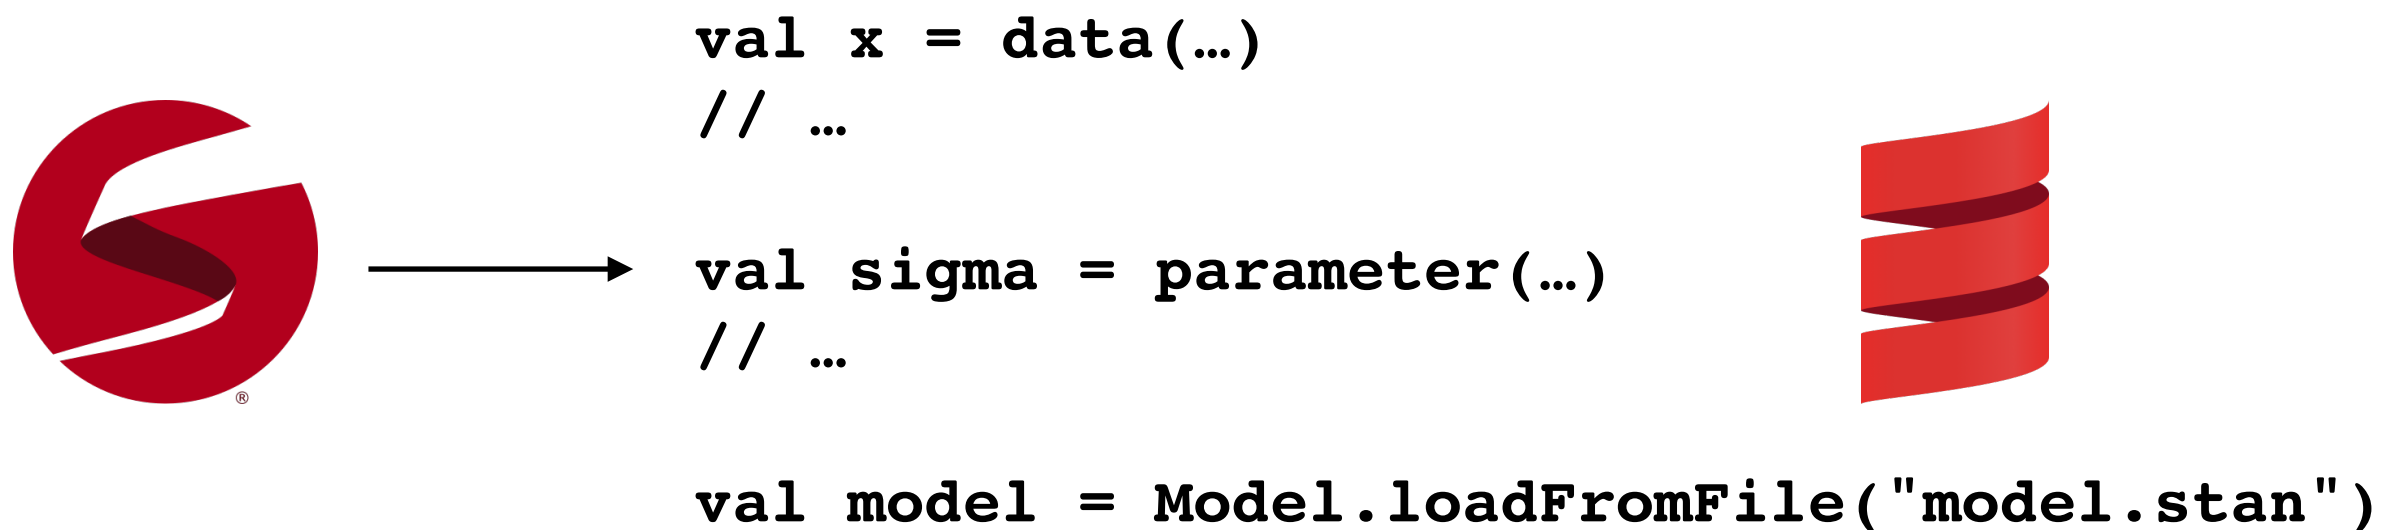


Integration into the Stan Ecosystem

Model Re-use: Export



Model Re-use: Import



Future Directions

Higher-level API

- Library of parameterizable models

Stan code analysis and transformations


- Automatic re-parameterization
- Automatic extraction of transformations
- Traditional compiler passes applied to Stan

Common subexpression elimination

Constant Folding

Strength Reduction

Dead code elimination



Questions?

ScalaStan GitHub Repository:
<https://github.com/cibotech/ScalaStan>

We're Hiring Engineers and Data Scientists
of All Levels and Backgrounds!
<https://www.cibotechnologies.com/about/>